




INFOSERVE  
>eurodata-Gruppe



# Introduction to Kubernetes



**JUG Saar** Meetup 37 | 23.08.2018

THOMAS DARIMONT | Software Architect | eurodata AG  
DR. PHILIPP WALTER | Leiter IT | INFOSERVE GmbH

# Agenda

JUG Saar Meetup 37 | 23.08.2018

- 18:00 Begrüßung und Einführung in Kubernetes
- 18:30 Demo 1: lokaler Kubernetes-Cluster mit Vagrant  
Demo 2: Cloud Native Spring Boot App  
Demo 3: Rolling Upgrades
- 19:15 Pizza!
- 19:45 Demo 4: HTTP Session Clustering mit Hazelcast  
Demo 5: Keycloak Cluster als Helm-Chart  
Demo 6: Kubernetes-Cluster auf VMs selbst einrichten
- 20:15 Frage- und Feedbackrunde
- 20:30 Ende der Veranstaltung

VORWORT

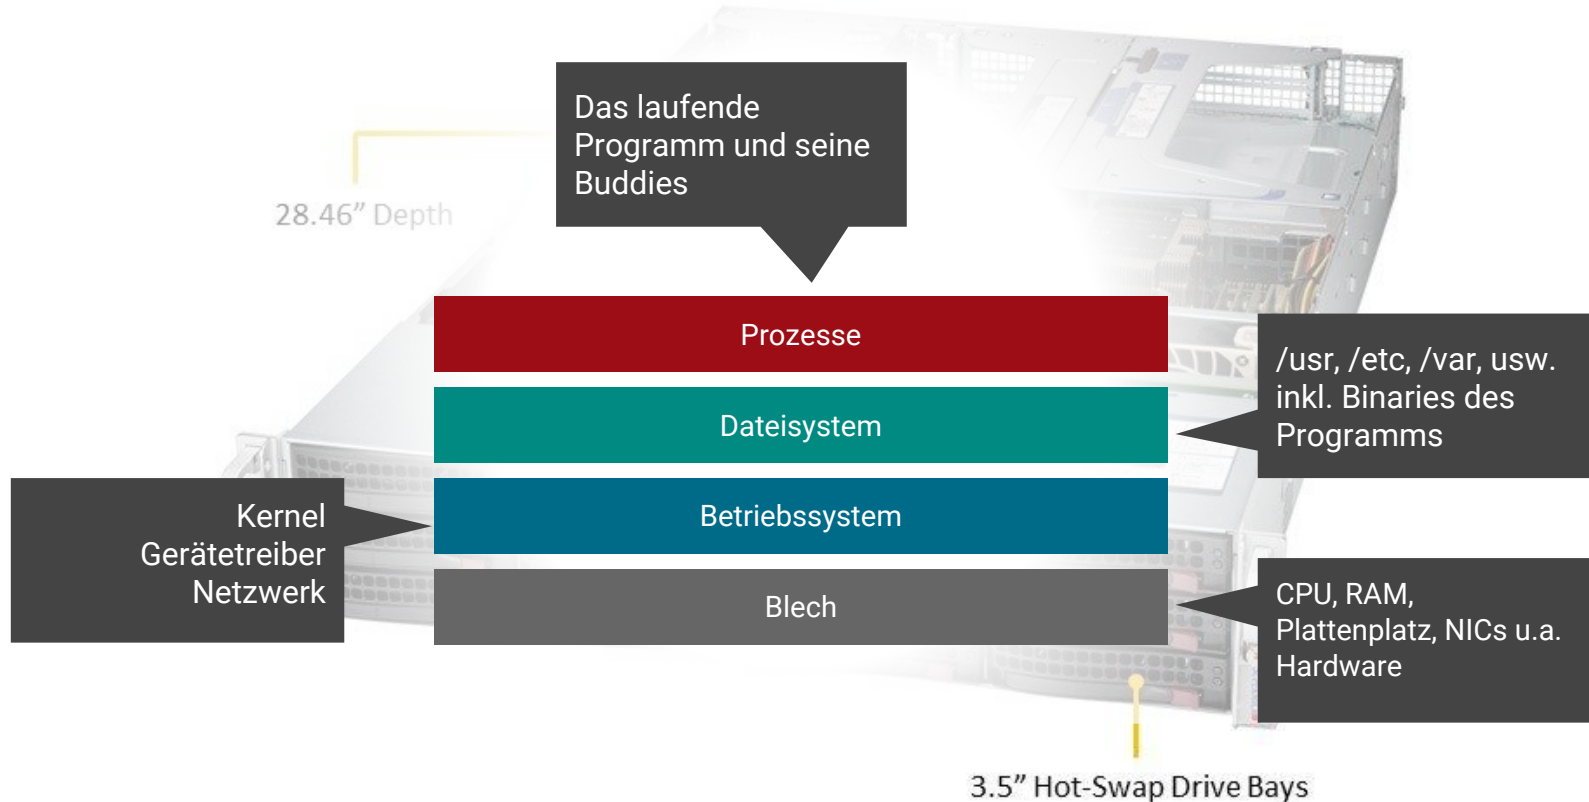
# Einführung in Kubernetes

DR. PHILIPP WALTER

<https://github.com/jugsaar/jugsaar-37-kubernetes-for-devs>

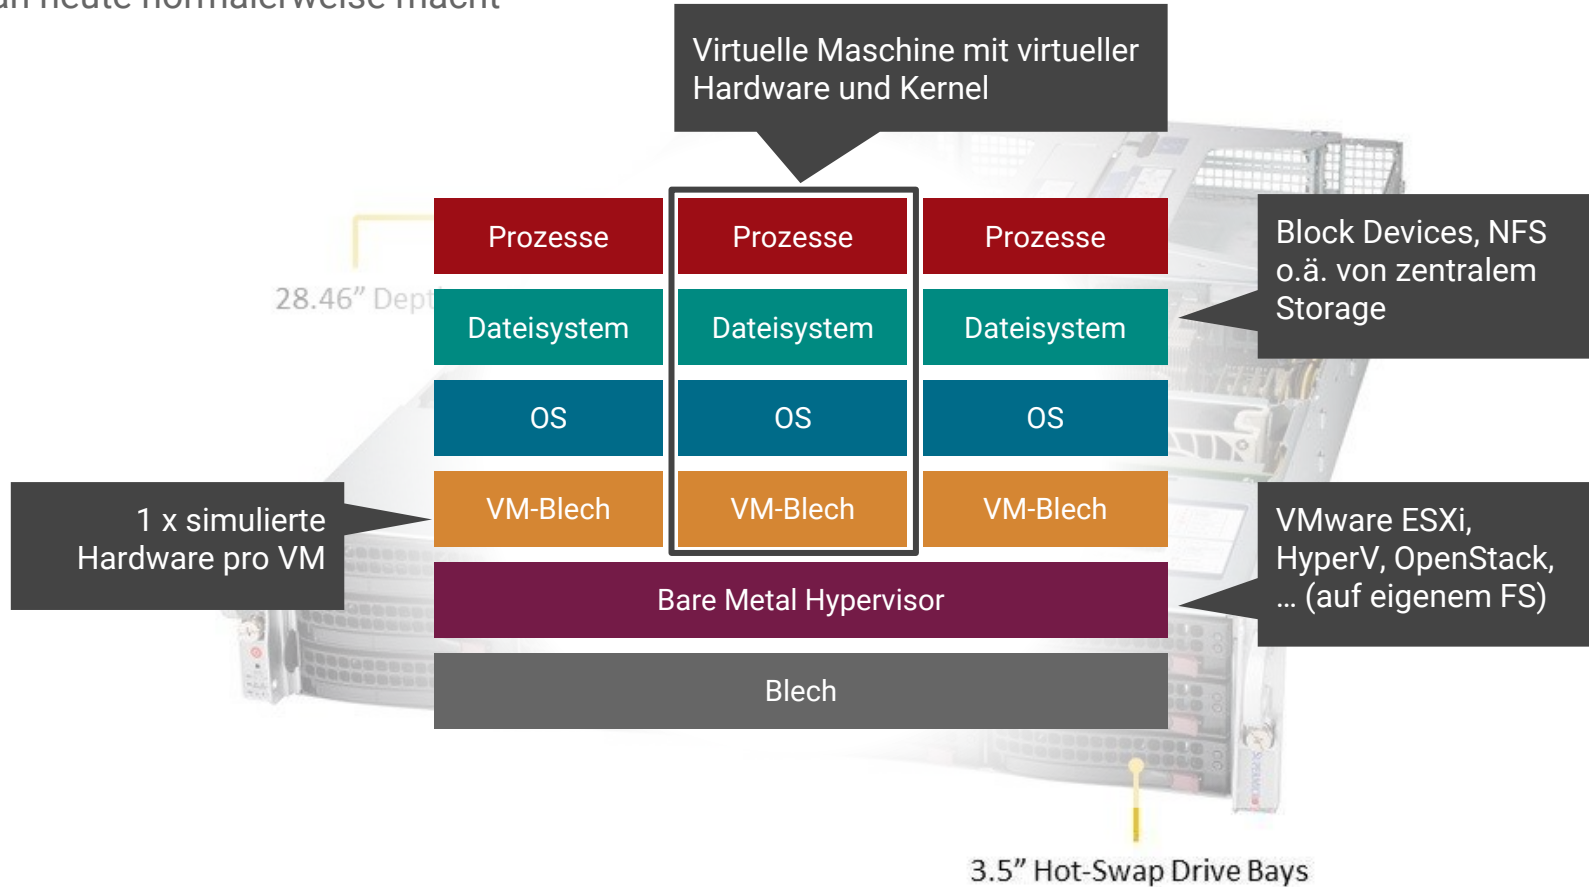
# Was braucht man, um ein Programm auszuführen?

Die reine Lehre



# Was braucht man, um ein Programm auszuführen?

Was man heute normalerweise macht



# Was braucht man, um ein Programm auszuführen?

Was man eigentlich bräuchte

Eine vertraute Umgebung, in der sich ein Prozess wohlfühlt. Könnte man auch einpacken und mitnehmen, um es anderswo laufen zu lassen...

Prozesse

Prozesse

Prozesse

Dateisystem

Dateisystem

Dateisystem

Container Runtime

Docker, rkt, OCI, ...

Betriebssystem


Linux mit namespaces und cgroups

Blech

3.5" Hot-Swap Drive Bays

# Was braucht man, um ein Programm auszuführen?

Container transportieren Software samt ihrer Ausführungsumgebung



```
epic ~ docker pull alpine
```

```
Using default tag: latest
```

```
latest: Pulling from library/alpine
```

```
8e3ba11ec2a2: Pull complete
```

```
Digest: sha256:7043076348bf5040220df6ad703798fd8593a0918d06d3ce30c6c93be117e430
```

```
Status: Downloaded newer image for alpine:latest
```

```
epic ~ docker run -it --rm alpine
```

```
/ #
```

# Container 101

Images suchen, herunterladen und starten

```
epic ~ docker search alpine
```

NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
alpine	A minimal Docker image based on Alpine Linux...	4092	[OK]	
mhart/alpine-node	Minimal Node.js built on Alpine Linux	373		
anapsix/alpine-java	Oracle Java 8 (and 7) with GLIBC 2.23 over A...	341		[OK]
...				

```
epic ~ docker pull alpine
```

```
Using default tag: latest
```

```
latest: Pulling from library/alpine
```

```
8e3ba11ec2a2: Pull complete
```

```
Digest: sha256:7043076348bf5040220df6ad703798fd8593a0918d06d3ce30c6c93be117e430
```

```
Status: Downloaded newer image for alpine:latest
```

```
epic ~ docker run -it --rm alpine
```

```
/ #
```



# Container 101

## Environment-Variablen von außen setzen

```
epic ~ docker run -it --rm -e MEINE_VARIABLE=MEIN_WERT alpine
```

```
/ # env
HOSTNAME=1529dea4b71d
SHLVL=1
HOME=/root
MEINE_VARIABLE=MEIN_WERT
TERM=xterm
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
PWD=/

/ # █
```

# Container 101

Host-Verzeichnisse und -Dateien als Container-Volumes mounten

```
epic ~ mkdir /tmp/demo
epic ~ echo huhu > /tmp/demo/hallo.txt
epic ~ docker run -it --rm -v /tmp/demo:/mnt alpine
/ # ls /mnt
hallo.txt
/ # cat /mnt/hallo.txt
huhu
```

Externe Shell

```
epic ~ echo Hallo Welt! > /tmp/demo/hallo.txt
```

```
/ # cat /mnt/hallo.txt
Hallo Welt!
/ #
```

# Container 101

Netzwerkdienste im Container von außen zugänglich machen

```
epic ~ docker run -it --rm -p 8080:80 alpine
```

```
/ # ip a
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
```

```
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```
    inet 127.0.0.1/8 scope host lo
```

```
        valid_lft forever preferred_lft forever
```

```
31: eth0@if32: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
```

```
    link/ether 02:42:ac:11:00:02 brd ff:ff:ff:ff:ff:ff
```

```
    inet 172.17.0.2/16 brd 172.17.255.255 scope global eth0
```

```
        valid_lft forever preferred_lft forever
```

```
/ # nc -l -p 80
```

Externe Shell

```
epic ~ echo huhu | nc -c localhost 8080
```

```
huhu
```

```
^C
```

```
/ # wget https://www.infoserve.de
```

```
Connecting to www.infoserve.de (212.89.154.36:443)
```

```
index.html          100% |*****| 65886    0:00:00 ETA
```

```
/ #
```

# Container 101

Registries, oder: wo kommen eigentlich die kleinen Container-Images her?

```
epic ~ docker run -d -p 5000:5000 --name registry registry:2
f8ebdc69117d2f927ab57ad43db063c9af5baac76a3340497a179e413a5199c3
```

```
epic ~ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
f8ebdc69117d	registry:2	"/entrypoint.sh /etc..."	3 seconds ago	Up 2 seconds	0.0.0.0:5000->5000/tcp	registry

```
epic ~ docker pull alpine
```

```
Status: Image is up to date for alpine:latest
```

```
epic ~ docker image tag alpine localhost:5000/mein_alpine
```

```
epic ~ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
alpine	latest	11cd0b38bc3c	6 weeks ago	4.41MB
localhost:5000/mein_alpine	latest	11cd0b38bc3c	6 weeks ago	4.41MB

```
epic ~ docker push localhost:5000/mein_alpine
```

```
The push refers to repository [localhost:5000/mein_alpine]
```

```
73046094a9b8: Pushed
```

```
latest: digest: sha256:0873c923e00e0fd2ba78041bfb64a105e1ecb7678916d1f7776311e45bf5634b size: 528
```

```
epic ~ docker run -it --rm localhost:5000/mein_alpine
```

```
/ #
```


# Kubernetes verwaltet Container auf Clustern

Docker : Kubernetes = 1 : N

	DOCKER	KUBERNETES
LAUFZEITUMGEBUNG	<b>Einzelner Host</b> mit Docker-Binaries	<b>Viele Hosts</b> mit Docker- und Kubernetes-Binaries
USE CASES	Software lokal laufen lassen Container-Images bauen und testen	Skalierbare und selbstheilende Produktionsumgebungen
API	docker pull, run, start, stop, ps, images, ...	kubectl + yaml-Specs (Details folgen...)
OBJEKTE	Container, Images, Volumes, Ports	Pods aus 1..n Containern, Services, Clusterknoten, Ingress, Deployments, ...
PARADIGMA	<b>Low-Level-Operationen</b> , die <b>manuell</b> auf einzelnen Containern ausgeführt werden	Definiere einen <b>Sollzustand</b> , den Kubernetes <b>automatisch</b> um jeden Preis zu halten versucht
SETUP	apt-get install docker	kubeadm, GKE, CoreOS, ... (Details folgen)

# Warum “Kubernetes”?


Ein bisschen Nerd muss sein: Google und die Borg



7 of 9, Borg

7 Seiten...

Interner Cluster Manager  
“Borg” bei Google (2015)



# kubernetes

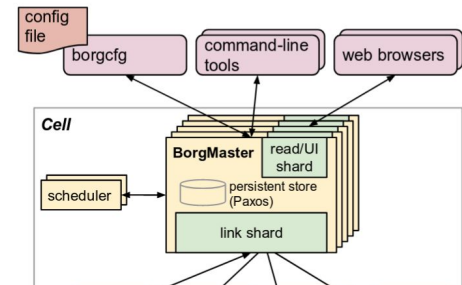
## Large-scale cluster management at Google with Borg

Abhishek Verma<sup>†</sup> Luis Pedrosa<sup>‡</sup> Madhukar Korupolu  
David Oppenheimer Eric Tune John Wilkes  
Google Inc.

### Abstract


Google’s Borg system is a cluster manager that runs hundreds of thousands of jobs, from many thousands of different applications, across a number of clusters each with up to tens of thousands of machines.

It achieves high utilization by combining admission control, efficient task-packing, over-commitment, and machine sharing with process-level performance isolation. It supports high-availability applications with runtime features that minimize fault-recovery time, and scheduling policies that re-



# Kubernetes 101

Laufzeitumgebung: Cluster statt Computer



```
epic ~ kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
kube01	Ready	master	3d	v1.11.2	<none>	CentOS Linux 7 (Core)	3.10.0-862.11.6.el7.x86_64	docker://1.13.1
kube02	Ready	<none>	3d	v1.11.2	<none>	CentOS Linux 7 (Core)	3.10.0-862.11.6.el7.x86_64	docker://1.13.1
kube03	Ready	<none>	3d	v1.11.2	<none>	CentOS Linux 7 (Core)	3.10.0-862.11.6.el7.x86_64	docker://1.13.1

# Kubernetes 101

Kubernetes versucht, den Sollzustand zu erhalten, indem Prozesse am Laufen gehalten werden

```
epic ~ kubectl run -i -t busybox --image=busybox:1.28
```

If you don't see a command prompt, try pressing enter.

```
/ # exit
```

Session ended, resume using 'kubectl attach busybox-7999f69f9d-68tv4 -c busybox -i -t' command when the pod is running

```
epic ~ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
busybox-7999f69f9d-68tv4	1/1	Running	1	6m

```
epic ~ kubectl attach busybox-7999f69f9d-68tv4 -c busybox -i -t
```

If you don't see a command prompt, try pressing enter.

```
/ # exit
```

Session ended, resume using 'kubectl attach busybox-7999f69f9d-68tv4 -c busybox -i -t' command when the pod is running



```
epic ~ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
busybox-7999f69f9d-68tv4	1/1	Running	2	7m



# Kubernetes 101

Pods bestehen aus Containern und sind der elementare Baustein von Kubernetes



```
epic ~ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
busybox-799f69f9d-68tv4	1/1	Running	2	7m

```
epic ~ kubectl attach busybox-799f69f9d-68tv4 -c busybox -i -t
```

If you don't see a command prompt, try pressing enter.

```
/ #
```

# Kubernetes 101

Replica Sets halten Pods am Laufen und werden als Deployments verwaltet

```
epic ~ kubectl describe pod busybox-7999f69f9d-68tv4
```

```
...
```

```
Status:           Running
```

```
IP:               10.116.0.3
```

```
Controlled By:    ReplicaSet/busybox-7999f69f9d
```

```
...
```

```
epic ~ kubectl describe ReplicaSet/busybox-7999f69f9d
```

```
...
```

```
Controlled By:    Deployment/busybox
```

```
Replicas:         1 current / 1 desired
```

```
Pods Status:      1 Running / 0 Waiting / 0 Succeeded / 0 Failed
```

```
...
```

```
epic ~ kubectl describe Deployment/busybox
```

```
...
```

```
Replicas:         1 desired | 1 updated | 1 total | 1 available | 0 unavailable
```

```
StrategyType:      RollingUpdate
```

```
RollingUpdateStrategy: 1 max unavailable, 1 max surge
```

```
...
```

# Kubernetes 101

## Deployments bündeln Pods und Replica Sets

```
epic ~ kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/busybox	1	1	1	1	1h

NAME	DESIRED	CURRENT	READY	AGE
rs/busybox-7999f69f9d	1	1	1	1h

NAME	READY	STATUS	RESTARTS	AGE
po/busybox-7999f69f9d-68tv4	1/1	Running	3	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3d

```
epic ~ kubectl delete deployment busybox
```

```
deployment "busybox" deleted
```

```
epic ~ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
po/busybox-7999f69f9d-68tv4	1/1	Terminating	3	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3d

# Deployments

Deployments können als YAML oder JSON beschrieben werden. Hurra, Infrastructure as Code!

echo.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: echo
spec:
  replicas: 2
  selector:
    matchLabels:
      pod: echo
  template:
    metadata:
      labels:
        pod: echo
    spec:
      containers:
        - name: echo
          image: jmalloc/echo-server:latest
          ports:
            - containerPort: 8080
```

```
epic ~ kubectl apply -f echo.yaml
```

```
deployment "echo" created
```

```
epic ~ kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/echo	2	2	2	2	6m

NAME	DESIRED	CURRENT	READY	AGE
rs/echo-8594ff85cb	2	2	2	6m

NAME	READY	STATUS	RESTARTS	AGE
po/echo-8594ff85cb-vg72l	1/1	Running	0	6m
po/echo-8594ff85cb-xbl9t	1/1	Running	0	6m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5d

```
epic ~ kubectl scale deployment echo --replicas 3
```

```
deployment "echo" scaled
```

```
epic ~ kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/echo	3	3	3	3	7m

```
...
```

# Exkurs: Labels und Selektoren

Man kann alles beschriften und damit auswählen

Selector: app=infos

app = infos  
stage = prod  
tier = web

app = infos  
stage = test  
tier = web

app = infos  
stage = prod  
tier = app

app = infos  
stage = test  
tier = app

Selector: app=ed, tier=web

app = ed  
stage = prod  
tier = web

app = ed  
stage = test  
tier = web

app = ed  
stage = prod  
tier = app

app = ed  
stage = test  
tier = app

Selector:  
tier=db

app = infos  
stage = prod  
tier = db


app = infos  
stage = test  
tier = db

app = ed  
stage = prod  
tier = db

app = ed  
stage = test  
tier = db

# Services

Services ermöglichen Layer-4-Load-Balancing über Pods



echo-service.yaml

```
kind: Service
apiVersion: v1
metadata:
  name: echo-service
spec:
  selector:
    pod: echo
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
    type: NodePort
```

```
epic ~ kubectl apply -f echo-service.yaml
service "echo-service" created
```

```
epic ~ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
echo-service	NodePort	10.104.0.183	<none>	80:31763/TCP	6m
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5d

# Services

Die Cluster-IP eines Service kann überall im Cluster auch über seinen DNS-Namen aufgelöst werden

```
epic ~ kubectl run -i -t busybox --image=busybox:1.28
```

If you don't see a command prompt, try pressing enter.

```
/ # nslookup echo-service
```

```
Server:      10.96.0.10
```

```
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local
```

```
Name:        echo-service
```

```
Address 1: 10.104.0.183 echo-service.default.svc.cluster.local
```

```
/ # watch -n 1 wget -qO - echo-service | grep served
```

```
Request served by echo-8594ff85cb-vg72l
```

```
Request served by echo-8594ff85cb-xbl9t
```

```
Request served by echo-8594ff85cb-xbl9t
```

```
Request served by echo-8594ff85cb-vg72l
```

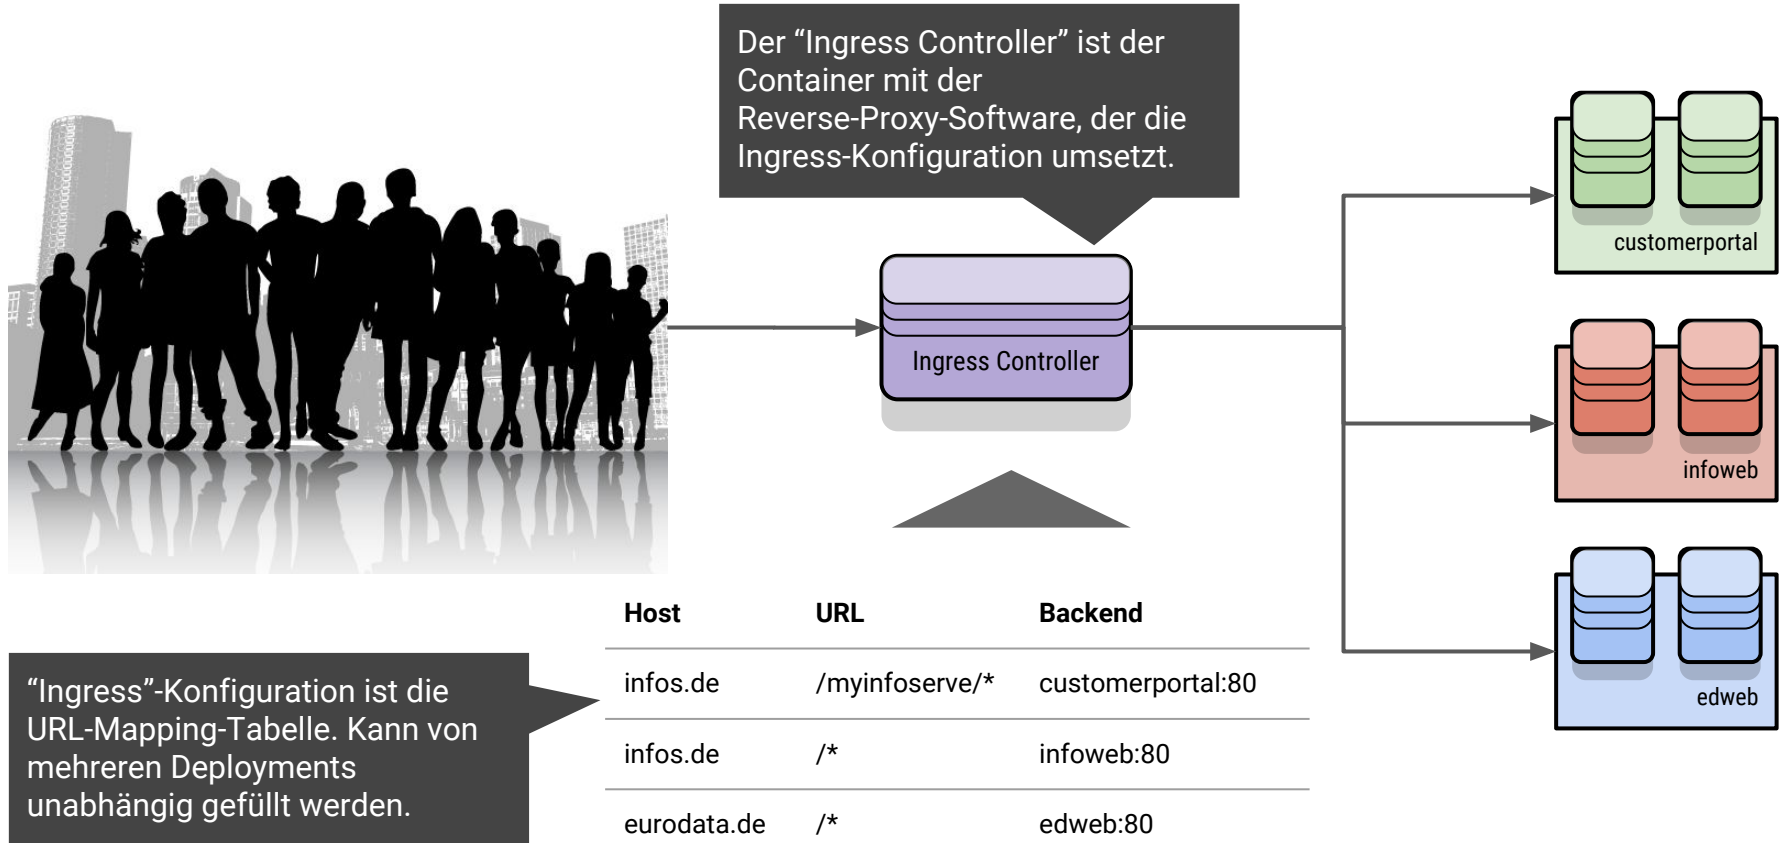
```
Request served by echo-8594ff85cb-vg72l
```

```
...
```

```
/ #
```

# Ingress Controller

Layer-7-Load-Balancing mit URL Mappings auf verschiedene Services






# Ingress Controller

Ingress Controller auf eigenem Cluster (kein GKE) deployen

traefik-ingress-controller.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: traefik-ingress-controller
  namespace: kube-system
---
apiVersion: extensions/v1beta1
kind: DaemonSet
metadata:
  ...
```



<https://docs.traefik.io/user-guide/kubernetes/>

```
epic ~ kubectl get all --namespace=kube-system -l k8s-app=traefik-ingress-lb
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
ds/traefik-ingress-controller-v1	2	2	2	2	2	<none>	8m

NAME	READY	STATUS	RESTARTS	AGE
po/traefik-ingress-controller-v1-7664s	1/1	Running	0	8m
po/traefik-ingress-controller-v1-pk7bq	1/1	Running	0	8m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
traefik-ingress-service	ClusterIP	10.99.243.201	<none>	80/TCP,443/TCP,8080/TCP	8m	k8s-app=traefik-ingress-lb

# Ingress-Ressourcen

Ingress-Ressourcen machen Services unter einer URL und einem Pfad zugänglich

echo-ingress.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: echo-ingress
spec:
  rules:
  - http:
      paths:
      - path: /echo
        backend:
          serviceName: echo-service
          servicePort: 80
```

```
epic ~ kubectl create -f echo-ingress.yaml
ingress "echo-ingress" created
```

```
epic ~ curl kube02/echo
Request served by echo-8594ff85cb-xmbb7
```

HTTP/1.1 GET /echo

```
Host: kube02
X-Forwarded-Port: 80
X-Forwarded-Server: traefik-ingress-controller-qvqbj
X-Real-IP: 10.124.0.0
User-Agent: curl/7.61.0
X-Forwarded-For: 10.124.0.0
X-Forwarded-Host: kube02
Accept-Encoding: gzip
Accept: */*
X-Forwarded-Proto: http
```

```
epic ~ curl kube03/echo
Request served by echo-8594ff85cb-mj65h
```

HTTP/1.1 GET /echo

...

DEMO 1

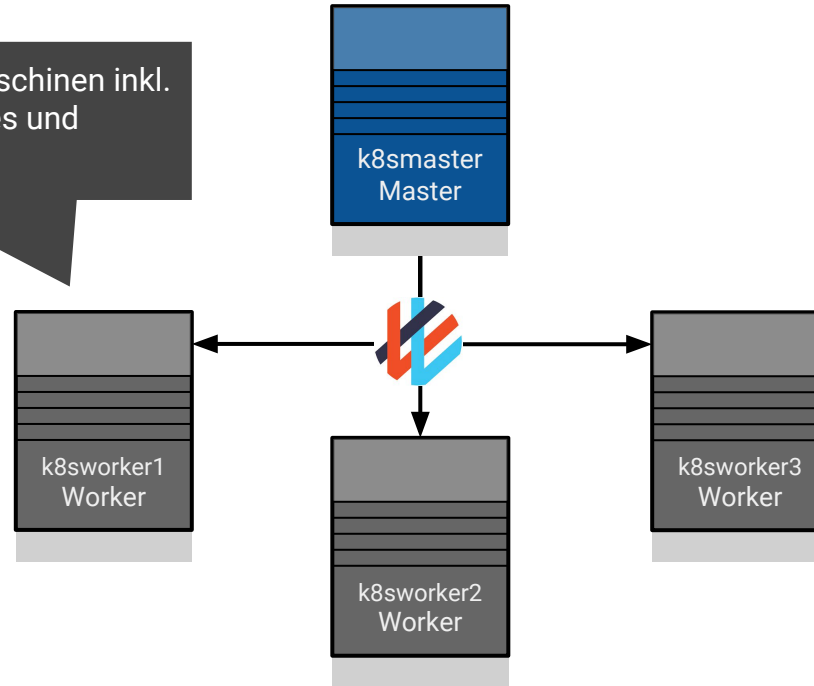
# Lokaler Kubernetes-Cluster mit Vagrant

THOMAS DARIMONT

<https://github.com/thomasdarimont/vagrant-kubernetes-lab/tree/poc/kubernetes-next>

# Lokaler Kubernetes-Cluster mit Vagrant

Vagrant richtet lokale virtuelle Maschinen inkl. Netzwerkkonfiguration, Kubernetes und WeaveNet als CNI-Plugin ein.



DEMO 2

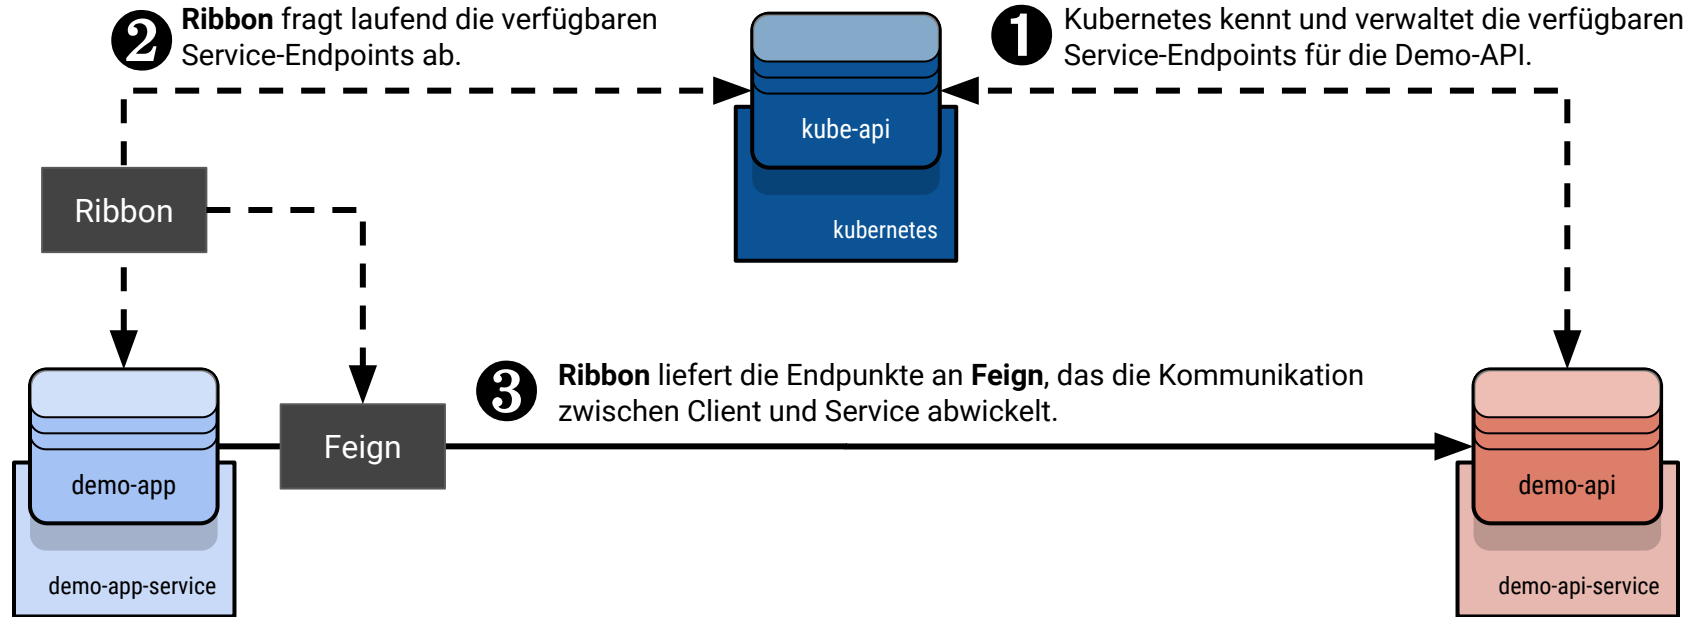
# Cloud-Native Spring Boot App

THOMAS DARIMONT

<https://github.com/jugsaar/jugsaar-37-kubernetes-for-devs/tree/master/code/spring-boot-k8s-demo>

# Cloud Native Spring Boot App

inkl. REST-Call mit Feign und Client-side Load Balancing mit Ribbon



DEMO 3

# Rolling Upgrade

THOMAS DARIMONT

<https://github.com/jugsaar/jugsaar-37-kubernetes-for-devs/tree/master/code/rolling-upgrade-demo>

DEMO 4

# HTTP Session Clustering mit Hazelcast

THOMAS DARIMONT

<https://github.com/jugsaar/jugsaar-37-kubernetes-for-devs/tree/master/code/spring-boot-k8s-hazelcast>

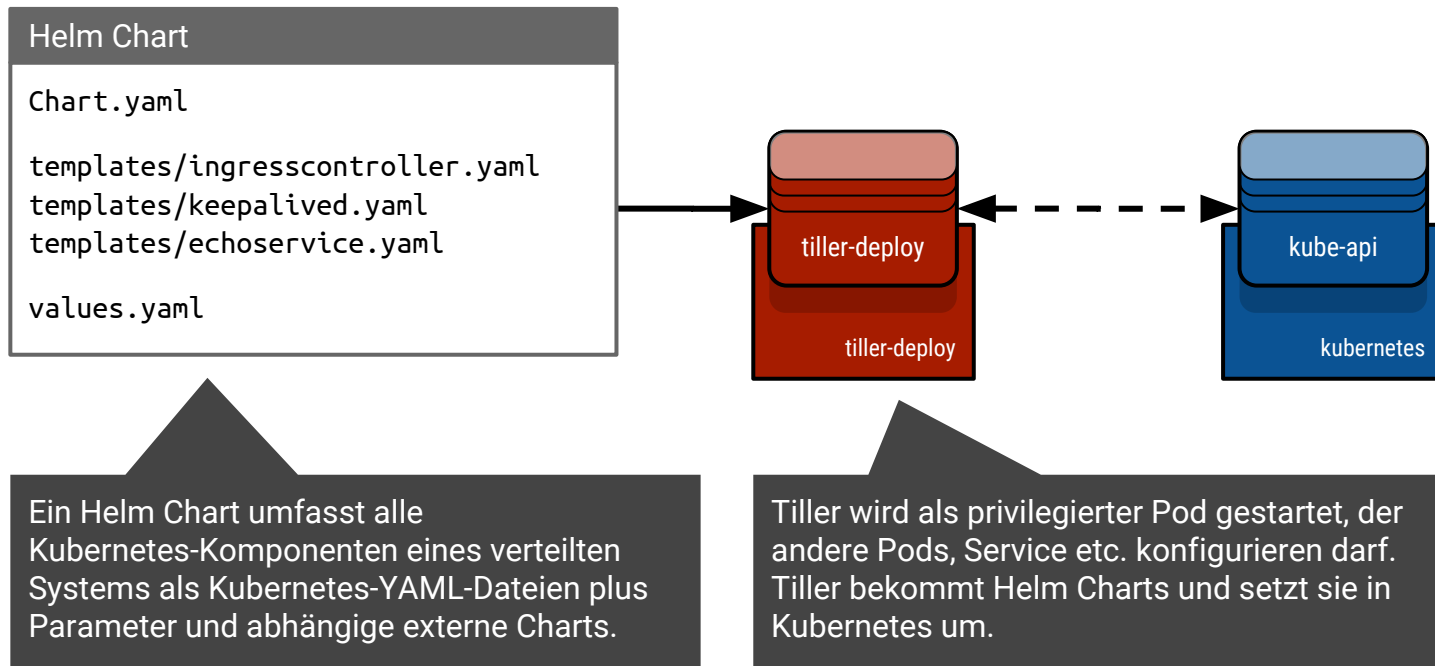


DEMO 5

# Keycloak Cluster als Helm Chart

THOMAS DARIMONT

<https://github.com/jugsaar/jugsaar-37-kubernetes-for-devs/tree/master/code/keycloak-helm>



DEMO 6

# Kubernetes-Cluster auf VMware-Basis

DR. PHILIPP WALTER

<https://github.com/jugsaar/jugsaar-37-kubernetes-for-devs>